

OGC CDB Vector Data in GeoPackage Interoperability Experiment

Table of Contents

1. Summary	4
1.1. Requirements & Research Motivation	4
1.2. Prior-After Comparison	4
1.3. Recommendations for Future Work	4
1.4. Document contributor contact points	5
1.5. Foreword	5
2. References	7
3. Terms and definitions	8
3.1. Abbreviated terms	8
4. Overview	9
4.1. IE Overview	9
4.2. Engineering Report Section Overviews.	9
5. Details of the Experiments Performed	10
5.1. A CDB Data Store - General Overview	10
5.2. Details on Experiment 1: CDB Data Used	10
5.3. The CDB Data Stores used in this IE	11
5.3.1. << >>	11
5.3.2. Presagis	11
5.3.3. VATC	11
5.4. Tools used to convert Shapefiles to GeoPackages	11
5.4.1. VATC	11
5.5. Key discussion topics related to CDB data stores	11
5.6. Details Related to Experiment 2	12
5.6.1. Participation in Experiment 2	12
5.6.2. Option 1a – 1:1 Conversion of Shapefiles to GeoPackages	12
5.6.3. Option 1b – Conversion of Shapefiles to GeoPackages using Normalized SQL Data	13
5.6.4. Option 1c – Flattened Attribution	14
5.6.5. Option 1d – Flattened Attribution + extensions	14
5.7. Experiment 3 - Each CDB LoD as a layer in GeoPackage	15
5.8. Experiment 4: Store each Geocell of Vector Data as a layer in GeoPackage	15
6. IE Experiment Results	17
6.1. Aechelon Technology IE Report	17
6.1.1. Use case and experiment focus	17
6.1.2. Aechelon Experiment Methodology	17
6.1.3. Metrics	18
6.1.4. Legend	20
6.1.5. Notes and observations	20
6.1.6. Compusult metrics from Experiment 2	21

6.2. FlightSafety Experiment Results	21
6.2.1. FlightSafety International's Use Case for CDB	21
6.2.2. FlightSafety Experiment Focus	22
6.2.3. FlightSafety Experiment Methodology	22
6.2.4. FlightSafety Metrics	23
6.2.5. Shapefile vs. GeoPackage Option 1 (Experiment 2) Testing	24
6.2.6. GeoPackage Option 3 and 4 Testing	27
6.2.7. Further GeoPackage Option 3 & 4 Testing	30
7. Guidance	33
7.1. Who is doing what	33
7.1.1. Hexagon/Luciad	33
7.2. Metrics captured	33
7.2.1. FSI	33
8. Recommendations, Observations, and Conclusions	34
8.1. Backwards compatibility in the CDB context	34
8.1.1. FlightSafety observations on backwards compatibility	34
8.2. Conclusions	34
8.2.1. Aechelon	34
8.2.2. FlightSafety	35
8.3. Recommendations	35
8.3.1. FlightSafety	35
Appendix A: Abstract Test Suite	37
A.1. Test module for conformance level 1	37
A.1.1. Conformance level 1	37
A.1.2. Test case for validity of XML response entity	37
A.2. Test module for conformance level 2	38
A.2.1. Conformance level 2	38
A.2.2. Test case for validity of XML response entity	38
Appendix B: Revision History	39
Appendix C: Bibliography	40

Publication Date: 2019-MM-DD

Approval Date: 2019-MM-DD

Submission Date: 2019-MM-DD

Reference number of this document: OGC 19-007

Reference URL for this document: <http://www.opengis.net/doc/PER/CDB-GPKG-IE>

Category: OGC Public Engineering Report

Editor: Carl Reed, PhD

Title: OGC CDB Vector Data in GeoPackage Interoperability Experiment

OGC Public Engineering Report

COPYRIGHT

Copyright © 2019 Open Geospatial Consortium. To obtain additional rights of use, visit <http://www.opengeospatial.org/>

WARNING

This document is not an OGC Standard. This document is an Engineering Report created as a deliverable in an OGC Interoperability Initiative and is not an official position of the OGC membership. This Engineering Report is distributed for review and comment. The content is subject to change without notice and may not be referred to as an OGC Standard. Further, any Engineering Report should not be referenced as required or mandatory technology in procurements. However, the discussions in this document could very well lead to the definition of an OGC Standard.

LICENSE AGREEMENT

Permission is hereby granted by the Open Geospatial Consortium, ("Licensor"), free of charge and subject to the terms set forth below, to any person obtaining a copy of this Intellectual Property and any associated documentation, to deal in the Intellectual Property without restriction (except as set forth below), including without limitation the rights to implement, use, copy, modify, merge, publish, distribute, and/or sublicense copies of the Intellectual Property, and to permit persons to whom the Intellectual Property is furnished to do so, provided that all copyright notices on the intellectual property are retained intact and that each person to whom the Intellectual Property is furnished agrees to the terms of this Agreement.

If you modify the Intellectual Property, all copies of the modified Intellectual Property must include, in addition to the above copyright notice, a notice that the Intellectual Property includes modifications that have not been approved or adopted by LICENSOR.

THIS LICENSE IS A COPYRIGHT LICENSE ONLY, AND DOES NOT CONVEY ANY RIGHTS UNDER ANY PATENTS THAT MAY BE IN FORCE ANYWHERE IN THE WORLD. THE INTELLECTUAL PROPERTY IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, AND NONINFRINGEMENT OF THIRD PARTY RIGHTS. THE COPYRIGHT HOLDER OR HOLDERS INCLUDED IN THIS NOTICE DO NOT WARRANT THAT THE FUNCTIONS CONTAINED IN THE INTELLECTUAL PROPERTY WILL MEET YOUR REQUIREMENTS OR THAT THE OPERATION OF THE INTELLECTUAL PROPERTY WILL BE UNINTERRUPTED OR ERROR FREE. ANY USE OF THE INTELLECTUAL PROPERTY SHALL BE MADE ENTIRELY AT THE USER'S OWN RISK. IN NO EVENT SHALL THE COPYRIGHT HOLDER OR ANY CONTRIBUTOR OF INTELLECTUAL PROPERTY RIGHTS TO THE INTELLECTUAL PROPERTY BE LIABLE FOR ANY CLAIM, OR ANY DIRECT, SPECIAL, INDIRECT OR CONSEQUENTIAL DAMAGES, OR ANY DAMAGES WHATSOEVER RESULTING FROM ANY ALLEGED INFRINGEMENT OR ANY LOSS OF USE, DATA OR PROFITS, WHETHER IN AN ACTION OF CONTRACT, NEGLIGENCE OR UNDER ANY OTHER LEGAL THEORY, ARISING OUT OF OR IN CONNECTION WITH THE IMPLEMENTATION, USE, COMMERCIALIZATION OR PERFORMANCE OF THIS INTELLECTUAL PROPERTY.

This license is effective until terminated. You may terminate it at any time by destroying the Intellectual Property together with all copies in any form. The license will also terminate if you fail to comply with any term or condition of this Agreement. Except as provided in the following sentence, no such termination of this license shall require the termination of any third party end-user sublicense to the Intellectual Property which is in force as of the date of notice of such termination. In addition, should the Intellectual Property, or the operation of the Intellectual Property, infringe, or in LICENSOR's sole opinion be likely to infringe, any patent, copyright, trademark or other right of a third party, you agree that LICENSOR, in its sole discretion, may terminate this license without any compensation or liability to you, your licensees or any other party. You agree upon termination of any kind to destroy or cause to be destroyed the Intellectual Property together with all copies in any form, whether held by you or by any third party.

Except as contained in this notice, the name of LICENSOR or of any other holder of a copyright in all or part of the Intellectual Property shall not be used in advertising or otherwise to promote the sale, use or other dealings in this Intellectual Property without prior written authorization of LICENSOR or such copyright holder. LICENSOR is and shall at all times be the sole entity that may authorize you or any third party to use certification marks, trademarks or other special designations to

indicate compliance with any LICENSOR standards or specifications.

This Agreement is governed by the laws of the Commonwealth of Massachusetts. The application to this Agreement of the United Nations Convention on Contracts for the International Sale of Goods is hereby expressly excluded. In the event any provision of this Agreement shall be deemed unenforceable, void or invalid, such provision shall be modified so as to make it valid and enforceable, and as so modified the entire Agreement shall remain in full force and effect. No decision, action or inaction by LICENSOR shall be construed to be a waiver of any rights or remedies available to it.

None of the Intellectual Property or underlying information or technology may be downloaded or otherwise exported or reexported in violation of U.S. export laws and regulations. In addition, you are responsible for complying with any local laws in your jurisdiction which may impact your right to import, export or use the Intellectual Property, and you represent that you have complied with any regulations or registration procedures required by applicable law to make this license enforceable.

Chapter 1. Summary

This OGC Engineering Report (ER) documents the results of the CDB Vector Data in GeoPackage Interoperability Experiment (IE). The participants in this IE tested the encoding of vector data into one or more GeoPackage(s) and storing the result in a CDB data store. GeoPackage Version 1.2 and CDB Version 1.1 and related Best Practices were the standards baseline used for this experiment. The IE builds on the work described in the OGC "CDB, Leveraging GeoPackage" Discussion Paper.

A primary objective of this IE was to agree and document possible change requests and/or best practices for storing vector data in a CDB data other than ShapeFiles. These suggested changes requests and/or best/practices would be used as the basis for SDB SWG discussions related to possible revisions to the CDB standard.

Once the IE activity is completed, Insert key findings here.

Just a note for now. Aechelon and FlightSafety - Shapefile \Rightarrow GeoPackage good for 1.2

1.1. Requirements & Research Motivation

In the commercial modelling and simulation industry, geospatial data content and use is exploding at a rate that is outpacing the innovation and utilization of the traditional M&S industry. The M&S and geospatial-intelligence (GEOINT) industries are on a path of convergence. Within the OGC, there are two geospatial standards that best enable the unification of the M&S and GEOINT industries: OGC CDB and GeoPackage. OGC CDB and GeoPackage are both standards increasingly used in M&S and GEOINT industry, but they both contain weaknesses and strengths when it comes to the combined needs of both industries.

1.2. Prior-After Comparison

This IE is based on preliminary work conducted under funding by the U.S. Army Geospatial Center (AGC), the results of which have been publicly released as OGC Document Number 18-077: OGC CDB and GeoPackage Discussion Paper (https://portal.opengeospatial.org/files/?artifact_id=80537&version=1). The Discussion Paper documents the results of research, design, and prototype efforts to present the OGC with an approach to creating "GeoCDB" — a technology mashing of GeoPackage and OGC CDB — as a deterministic repository of easily read data geospatial datasets suitable for storage, runtime access, and dissemination for live, virtual, constructive, gaming, and mission command (MC) systems.

1.3. Recommendations for Future Work

This Interoperability Experiment addressed the specific issue of how best to and develop guidance for storing GeoPackage containers in a CDB datastore. AS a result of the experiments conducted in this IE activity, the particants have identified a number of other potential interoperability experiments that could be conducted in the future.

This current and potential future IEs will provide highly relevant input into the specification and development of CDB Version 2.0.

1.4. Document contributor contact points

All questions regarding this document should be directed to the editor or the contributors:

Contacts

Role	Name	Organization
Initiator/Lead	David Graham	CAE
Technical Lead	Carl Reed	Carl Reed & Associates
editor	Carl Reed	Carl Reed & Associates
Initiator	Kevin Bentley	Cognitics
Initiator	Holly Black	CACI
Initiator	Hermann Brassard	Presagis
Initiator	Brian Ford	FlightSafety
Initiator	Ryan Franz	FlightSafety
Initiator	Jay Freeman	CAE
Initiator	Glen Johnson	VATC
Initiator	Mike Lokuta	CAE
Initiator	Bernard Leclerc	CAE
Initiator	Lance Marrou	Leidos
Initiator	Earl Miller	USSOCOM
Initiator	Ronald Moore	Leidos
Initiator	Susan Raymie	USSOCOM
Initiator	Sara Saeedi	U of Calgary
Participant	Jordan Dauble	Simblocks.io
Participant	Paul Foley	KaDSci
Participant	Priamos Georgiades	Aechelon
Participant	Robert Goff	Hexagon US Federal
Participant	Dave Lajoie	Presgis
Participant	Graham Long	Thales
Participant	Jason McDonald	Compusult
Participant	Nacho Sans-Pastor	Aechelon
Participant	Kathleen Ski	ISPA Technology
Participant	Trent Tinker	Luciad

1.5. Foreword

Attention is drawn to the possibility that some of the elements of this document may be the subject of patent rights. The Open Geospatial Consortium shall not be held responsible for identifying any or all such patent rights.

Recipients of this document are requested to submit, with their comments, notification of any relevant patent claims or other intellectual property rights of which they may be aware that might be infringed by any implementation of the standard set forth in this document, and to provide supporting documentation.

Chapter 2. References

The following normative documents are referenced in this Engineering Report.

- [OGC 15-113r5](https://portal.opengeospatial.org/files/15-113r5) [https://portal.opengeospatial.org/files/15-113r5], Volume 1: OGC CDB Core Standard: Model and Physical Data Store Structure
- [OGC 16-070r3](https://portal.opengeospatial.org/files/16-070r3) [https://portal.opengeospatial.org/files/16-070r3], Volume 4: OGC CDB Best Practice use of Shapefiles for Vector Data Storage

The following informative documents are referenced in this Engineering Report.

- [OGC 18-077r2, OGC CDB, Leveraging GeoPackage Discussion Paper](https://portal.opengeospatial.org/files/?artifact_id=82553) [https://portal.opengeospatial.org/files/?artifact_id=82553]

Chapter 3. Terms and definitions

For the purposes of this report, the definitions specified in Clause 4 of the OWS Common Implementation Standard [OGC 06-121r9](https://portal.opengeospatial.org/files/?artifact_id=38867&version=2) [https://portal.opengeospatial.org/files/?artifact_id=38867&version=2] shall apply. In addition, the following terms and definitions apply.

- term name

text of the definition

- term name|synonym

text of the definition

3.1. Abbreviated terms

The following abbreviated terms provides a list of the abbreviated terms and the symbols necessary for understanding this document.

CRS	Coordinate Reference System
GPKG	GeoPackage
ER	Engineering Report
IE	Interoperability Experiment
IG	Image Generation
OGC	Open Geospatial Consortium
OGR	OGR Simple Features Library
OWS	OGC Web Services
SRS	Spatial Reference System
SWG	Standards Working Group

Chapter 4. Overview

4.1. IE Overview

The current approved OGC CDB 1.1 Standard provides a Best Practices document describing the encoding of vector data using Esri Shapefiles. All previous versions of CDB (including OGC CDB Standard 1.0 and the Industry-maintained CDB Specification 3.2 and prior) have only described vector data encoding in Shapefiles. All known CDB content and repositories that include vector data are encoded using Shapefiles. As a result, there have been numerous requests to test and document how other vector encodings could be stored in a CDB data store. This requirement is captured in OGC Change Request [545](http://ogc.standardstracker.org/show_request.cgi?id=545) [http://ogc.standardstracker.org/show_request.cgi?id=545]. Another identified issue is that every use of Shapefiles creates four physical files on storage which has performance implications.

The goal of this IE will be achieved by replicating the experimentation conducted in the Discussion Paper by IE participants who are CDB content creators, maintainers of CDB repositories, implementers of CDB datasets in run-time and non-run-time CDB use cases, or new CDB 1.1 users who are evaluating CDB implementations.

These IE participants will help produce an OGC Engineering Report that discusses whether the recommended alternate encoding of CDB vector data as a GeoPackage is fit for use and submission to the OGC for consideration as an approved change to the OGC CDB 1.1 Standard and/or Best Practices.

4.2. Engineering Report Section Overviews.

Section 5 introduces the background and requirements for adding GeoPackages as a storage format in a CDB data store. It describes the situation prior to the Interoperability Experiment and discusses the requirements defined by the CDB user base.

Section 6 discusses the experiments in detail. The concept of operation also provides recommendations on preferred strategies.

Section 7 discusses the results of the various experiments performed as part of this initiative. A clear mapping of requirements for proposed GeoPackages in a CDB datastore candidate standard/best practice is provided.

Section 8 provides a summary of the main findings and discusses proposed change requests.

Annex A provides detailed performance information.

Chapter 5. Details of the Experiments Performed

5.1. A CDB Data Store - General Overview

A CDB structured data store is physically arranged on disk into the following top level directory structures:

- Metadata contains a set of XML metadata and controlled vocabulary files that are global to the data store.
- GTModel contains geotypical models, generic models that are defined once in the CDB and are intended to be rendered in multiple places throughout the data store (contrast with geospecific or GS models). They aren't intended to represent specific objects, but simply a typical representation of an object type such as a tree.
- MModel contains Moving Models, which don't have a fixed location and are intended to be dynamically placed and moved throughout a simulation. An example is an automobile or aircraft.
- Tiles contains tiled datasets. This is the structure in which the majority of vector based Shapefiles are stored.
- Navigation contains global navigation datasets.

A CDB data stored is structured into partitions that organizes the world into 1-degree by 1-degree cells. Within each cell, the tiled datasets are organized into a Level of Detail (LOD) hierarchy. More specifically, the CDB storage model relies on three important means to organize the data: a) Tiles which organize the data into zones defined by its location with respect to a WGS84 reference system; b) Layers (or datasets) which organize different types of data in a tile; and c) Levels of Detail (LoD) which organize the data in each layer of each tile by its detail.

Most CDB datasets are organized in a tile structure and stored in the \CDB\Tiles\ directory. The tile structure facilitates access to the information in real-time by any runtime client devices. However, for some datasets that require minimal storage, such as, Moving Models or Geotypical Models, there is no significant advantage to be added for such a tile structure. Such datasets are referred to as global datasets and consist of data elements that are global to the earth.

Point, line, and area features are encoded in a CDB store as Shapefiles and are organized into several Vector Datasets and LoDs. For each LoD, the CDB standard specifies the maximum number of points allowed per Tile-LoD and the resulting average Feature Density is defined. All vector data in a CDB store are referenced to the WGS 84 Earth-centered, Earth-fixed terrestrial reference system and geodetic datum.

5.2. Details on Experiment 1: CDB Data Used

Three complete CDB data stores were provided for use in this Interoperability Experiment. The provision of CDB data stores and the subsequent conversion of Shapefiles into GeoPackages was defined in the IE Activity Plan in Experiment 1. The participants identified a number of activities

associated with executing Experiment 1.

1. Identify and document any differences in the Shapefile to GeoPackage conversions that resulted from the use of different conversion tools. .

5.3. The CDB Data Stores used in this IE

5.3.1. << >>

5.3.2. Presagis

The Presagis provided data store is for the country of Yemen.

5.3.3. VATC

The VATC provided CDB data store was created from fully open data sources. The data provided in this CDB data store are in one CDB Geocell containing downtown Los Angeles. The data include:

- USGS NAIP 1m background (Entire Geocell) (CDB Lod 7)
- USGS HighRes Ortho Program 1ft Ortho (Southern section of Geocell) (CDB Lod 9)
- USGS 1/3 NED Elevation (CDB Lod 4 for an entire geocell)
- OpenStreetMap Vector Map (Entire geocell)
- Los Angeles County Building footprint information. (Approximately 1,734,043 buildings were extruded from the footprint data)

For the IE, VATC made the data available at https://storage.cloud.google.com/epic_builder/OGC_IE/LosAngeles_CDB.zip?_ga=2.3746352.-1225582785.1543877247

5.4. Tools used to convert Shapefiles to GeoPackages

5.4.1. VATC

We will have link's to download our core 3rd Party as built libraries for windows. (linux users will need to build with cmake... depending on support library) and Github links for opensource lib's that have been modified to support CDB (this includes OpenSceneGraph, osgEarth, and GDAL) (note we do not anticipate any further fixes to the GDAL library will be required for this IE and the built GDAL version from the 3rd Party is expected to suffice for this experiment)

5.5. Key discussion topics related to CDB data stores

Given the size of the CDB data stores used in this IE, there was discussion related to how best to provide data stores updated with GeoPackage content. The general consensus was that downloading the entire data store just to get the GeoPackages was non-optimal and time wasteful. Therefore, the participants discussed using **Version** metadata (Volume 1 CDB Standard, Clause 5.1.8). They determined that the original CDB data with Shapefiles would be Version 1 and that CDB enhanced with GeoPackages would be Version 2.

5.6. Details Related to Experiment 2

Experiment 2 focused on approaches to replacing each Shapefile in an existing CDB data store thereby consolidating the three geometry files into a single GeoPackage. The objectives was to determine the best practices for not only replacing Shapefiles for also allowing the storage and use of both Shapefiles and Geopackages in a CDB data store. Part of this experiment was to also evaluate and compare performance using the baseline CDB datasets made available as part of Experiment 1. finally, this experiment also focused on evaluating and analyzing and results from Experiment #2 related to performance, backwards-compatibility and risks to interoperability.

5.6.1. Participation in Experiment 2

The following IE participants executed the following options defined for Experiment 2. Note: Details of the options are provided below.

- Aechelon - Options 1a, 1b & 1c with the Yemen and Los Angeles DBs
- CAE Montreal - option 1a with the numerous CDB databases that are available and perform a before-after comparison of the overall performance of the system.
- Compusult - All options
- FlightSafety - FlightSafety intends to test all the option 1 variants, as well as option 3 and 4. We are going to have more of a client focus with our testing, since that is more performance sensitive.
- Luciad - Option 1d
- Presagis - If time permit, we will also test the dynamic creation of Vector data when the simulation modifies the database at runtime.
- VATC - Supported Option 1b initially both with data creation and opensource visualization of this option. We will make a best effort to include visualization for Compusult's LA updates as well as our own (and others as published) . It is our intent as well to support one other option (option 3 most likely)

The Participants identified four possible approaches to converting and/or using GeoPackages in a CDB data store.

5.6.2. Option 1a – 1:1 Conversion of Shapefiles to GeoPackages

This experiment researched the direct 1 to 1 conversion of Shapefiles in a CDB datastore into a corresponding set of Geopackages. GDAL (https://www.gdal.org/drv_geopackage.html) and << >> were used to do the conversion. Charcateristics of the Option 1a approach are:

- There is a 4:1 reduction in the number of files.
- There is one layer (table) per GeoPackage.
- The Feature Class and Extended Attribute files have no geometry.
- “Off the Shelf” GeoPackage Viewers will have no compatibility over the feature class and extended attributes layers.
- This approach under-utilizes the capabilities of GeoPackage.




Figure 1. One to one conversion of Shapefiles to GeoPackages

5.6.3. Option 1b – Conversion of Shapefiles to GeoPackages using Normalized SQL Data

This experiment researched the approach of using normalized SQL in the conversion of Shapefiles into GeoPackages. This approach has the following characteristics:

- Utilizes a standard normalized relational database design, utilizing foreign keys.
- There is a 12:1 reduction in the number of files.
- There are three layers per GeoPackage.
- The Feature Class and Extended Attribute tables have no geometry.
- However, “Off the Shelf” GeoPackage software will not be aware of the extended and feature class attributes. This can be somewhat mitigated when a SQL View is used, which gives viewers (clients) read-only visibility over these attributes.

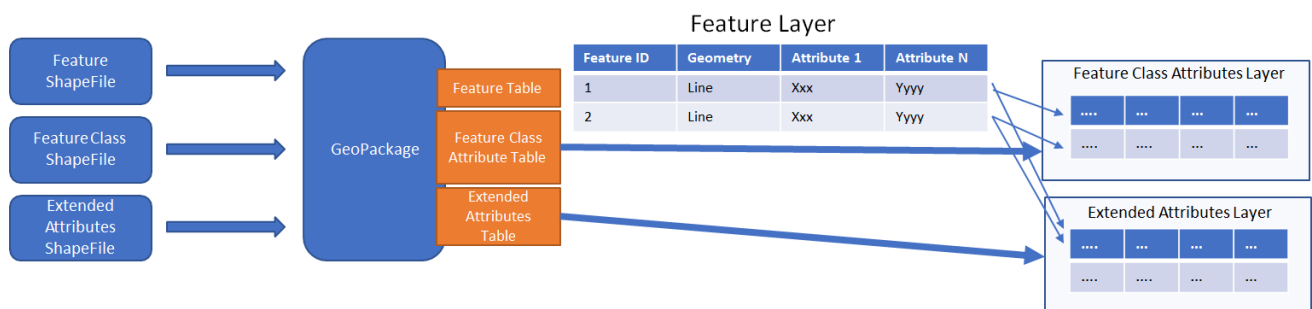


Figure 2. Use of Normalized SQL

5.6.4. Option 1c – Flattened Attribution

This experiment researched the approach of using flattened attribution in the conversion of Shapefiles into GeoPackages. This approach has the following characteristics:

- There is a 12:1 reduction in the number of files.
- Some duplication of data, resulting in larger files.
- There is one layer per GeoPackage.
- The Feature Class and Extended Attributes are populated for each feature.
- This approach utilizes a standard normalized relational database design, utilizing foreign keys.
- Full “Off the Shelf” GeoPackage software compatibility.

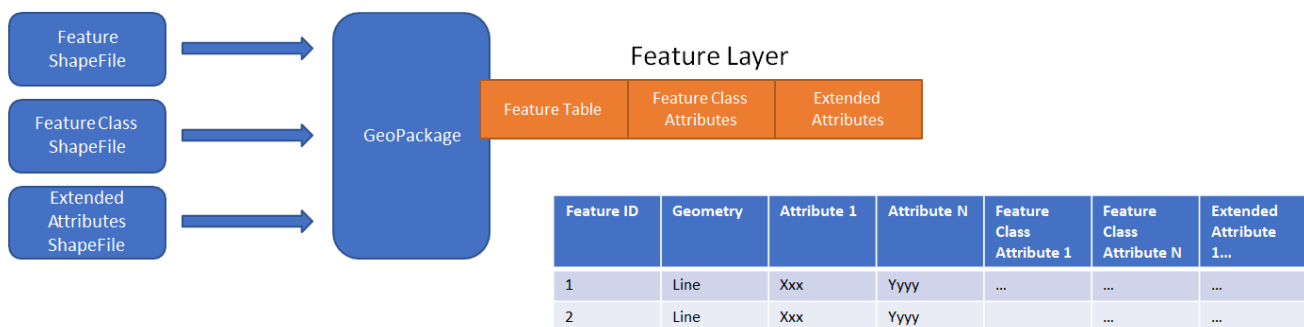


Figure 3. Flattened Attributes Approach

5.6.5. Option 1d – Flattened Attribution + extensions

This experiment researched the approach of using flattened attribution plus the GeoPackage related Tables extension in the conversion of Shapefiles into GeoPackages. The reason for using the extension was to enhance the ease of moving data in both directions (Shapefile to GeoPackage and visa-versa) using existing tools and without any data loss. This approach has the following characteristics:

- Flatten CDB standard instance and class attribute – maximum GIS tools compatibility
- “Off the Shelf” GeoPackage software compatibility for CDB standard attributes.
- Table (related tables) for extended attributes
- This approach utilizes a standard normalized relational database design, utilizing foreign keys.
- Some duplication of data, resulting in larger files (Class attributes).
- There is one layer per GeoPackage.
- The Feature Class and Extended Attributes are populated for each feature.

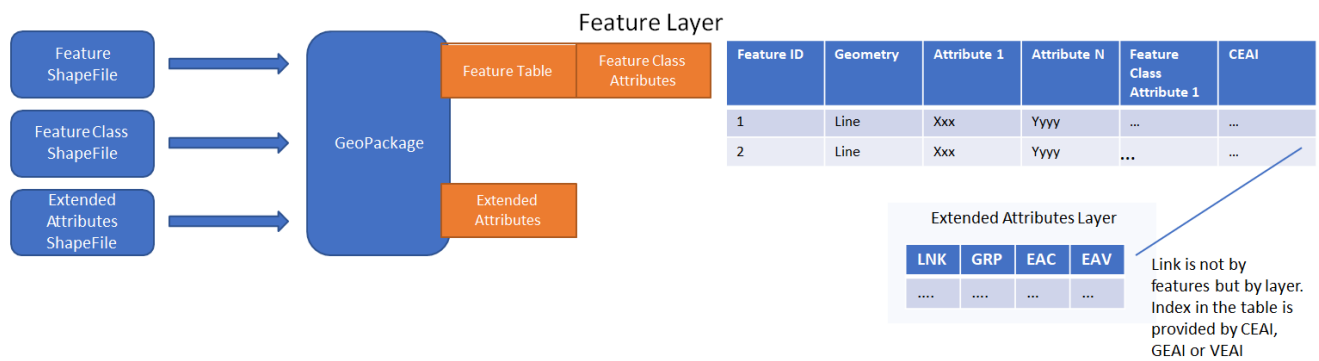


Figure 4. Flattened Attributes Approach

5.7. Experiment 3 - Each CDB LoD as a layer in GeoPackage

The methodology for Experiment 3 involves:

The goal of this experiment is to significantly reduce the number of files in both a CDB datastore and in the resulting GeoPackage. Steps in this experiment include:

- Modify implementation software to support storing an entire CDB Level of Detail (LoD) in a single GeoPackage.
- Evaluate and compare performance using the baseline CDB datasets and the Alternative #2 datasets.
- Evaluate analysis and results from Experiment #3 for performance, backwards-compatibility and risks to interoperability.

In this approach, the tables in the GeoPackage correspond to each LOD of CDB. The GeoPackage would contain 24 tables for each of the CDB LODs. Each CDB geotile would contain a GeoPackage to correspond to the CDB data stores (such road networks, geospecific points, etc.). CDB tiles for a data store combine into a single GeoPackage table within that given LOD where the tile definition (row and column) would be queryable attributes for each feature. In simple language, to find the features in a tile for a particular geotile's road network in LOD 3 of CDB, a consumer would open the road network GeoPackage, open the table that corresponds to LOD, and query for results where the column and row reference matches the CDB tile.

FlightSafety and Luciad are planning on running experiment 3.

5.8. Experiment 4: Store each Geocell of Vector Data as a layer in GeoPackage

The methodology involves:

This experiment extends Experiment 3 (above) to have a single GeoPackage per Geocell in a CDB datastore. This results in all LODs and all CDB feature layers in a single GeoPackage. The steps in this experiment include:

- Modify implementation software to support storing an entire GeoCell in a GeoPackage.
- Evaluate and compare performance using the baseline CDB datasets and the Alternative #3 datasets.
- Evaluate analysis and results from Experiment #4 for performance, backwards-compatibility and risks to interoperability.

In this approach, the tables in the GeoPackage correspond to each data store of CDB (such road networks, geospecific points, etc.). The GeoPackage would contain eight (8) layers representing each of the CDB data stores (GSFeature, GTFeature, GeoPolitical, VectorMaterial, RoadNetwork, RailRoadNetwork, PowerLineNetwork, and HydrographyNetwork). CDB tiles and LODs for a data store combine into a single GeoPackage table where the tile definition (row and column) and LOD would be queryable attributes for each feature. In simple language, to find the features in a location for a particular geotile's road network in LOD 3 of CDB, a consumer would open the geotile's GeoPackage, open the table that corresponds to data store, and query for results where the LOD column and row reference matches the CDB tile and LOD.

FlightSafety and Luciad are planning on running experiment 4.

Chapter 6. IE Experiment Results

This section of the Engineering Report provides details of the results of the experiments performed by each of the IE participants.

6.1. Aechelon Technology IE Report

6.1.1. Use case and experiment focus

The following is a description of a key use case in the Aechelon content processing workflow to use in an image generator.

- A CDB data store is used as the source for content to feed a publishing process into the Aechelon Image Generator (IG) runtime format. While reducing CDB storage requirements is a consideration, the primary concern in this workflow is with read access speed. Even so, the time taken by the 'feature scan' step of the publishing process, where the CDB source vector files are scanned to identify the features to import, is 2 orders of magnitude smaller than the rest of the pipeline. However, any improvements in speed and/or storage requirements have a positive impact on the efficiency of the Aechelon publication workflow.
- Note: After the feature scan step in the publishing process, all references to all features are in various python data structures, which are then given to the first of multiple processing steps to begin the data transformations. For example, point features with models will have their OpenFlight models converted to an intermediate Aechelon-specific model format, while their instance geographical data are saved in an Aechelon-specific lookup table format.

The publishing software is in python, with invocations of C++ EXEs for performance-critical processing. The feature scan step is entirely in Python, version 3.5. Please note that for the image generation workflow, only metadata fields that affect the appearance of features are considered and remainder of the CDB content is ignored, such as tactical data, or the entire geopolitical dataset. The hydrography network dataset is also ignored since the RMTexture dataset is used to identify areas of water.

For the purposes of this experiment, five feature types were considered and processed: Cultural, Lights, Powerlines, Railroads, and Trees. All of the performance information provided in the tables below are related to these five feature types.

The changes implemented on the publisher side to support GeoPackage were the minimal necessary to get functional parity with the Shapefile based implementation. In other words, no attempt was made to optimize the code to take advantage of the internals of the GeoPackage files using sqlite, and all data access went through the OGR module.

6.1.2. Aechelon Experiment Methodology

The following is a concise description of the methodology used for execution of the Aechelon committed tasks in this Interoperability Experiment.

- Following generation of the GeoPackage files for each option, changes were made as needed in the publishing scripts to be able to read and publish the data.

- Each feature type was tested by spot-checking in the image generator using a small reference database with representative data for each of the five feature types.
- The next step was to convert the following three of the four CDB data stores made available for the interoperability experiments. However, the entire publication end-to-end workflow for image generation was not performed due to the considerable time it would take for each run:
 - Yemen (4 geocells), from Presagis.
 - Downtown Los Angeles (1 geocell), from VATC.
 - Greater Los Angeles (4 geocells), from Cognitics.
- To generate data for Experiment 2, option 1A:
 - Ran Option1.py from the Cognitics conversion scripts in the master branch (<https://github.com/Cognitics/GeoCDB/tree/master>).
 - Deleted existing .shp, .shx, sidecar .dbf, .dbt & .prj files (i.e. kept .dbf files holding class/extended data.)
- To generate data for Experiment 2, option 1C:
 - Ran Option1.py from the Cognitics conversion scripts in the Presagis branch (<https://github.com/Cognitics/GeoCDB/tree/Presagis>).
 - Deleted the existing .shp, .shx, .dbf, .dbt & .prj files.
- To generate data for Experiment 2, option 1D:
 - Ran Option1d.py from the Cognitics conversion scripts in the master branch (after update of March 17, 2019, and some local edits to protect against 'None' during conversion of the LA databases.)
 - Deleted the existing .shp, .shx, .dbf, .dbt & .prj files.
- To generate data for Experiment 3:
 - Ran Option3.py from the Cognitics conversion scripts in the master branch (after update of March 17, 2019, and some local edits to uncomment writing the class metadata to the instance tables and to protect against 'None' during conversion of the LA databases.)
 - Deleted the 100_GSFeature, 101_GTFeature, 202_RailroadNetwork and 203_PowerlineNetwork folders from each geocell.
- To generate data for Experiment 4:
 - Ran Option4.py from the Cognitics conversion scripts in the master repository (after update of March 17, 2019, and some local edits to uncomment writing the class metadata to the instance tables and to protect against 'None' during conversion of the LA databases.)
 - Deleted the 100_GSFeature, 101_GTFeature, 202_RailroadNetwork and 203_PowerlineNetwork folders from each geocell.
- Then, for each option, disabled the publishing process beyond the 'feature scan' step and captured the following metrics for the three CDB data stores.

6.1.3. Metrics

The following three tables provide basic performance metrics for the three CDB data stores

processed in this experiment. Providing performance metrics is one of the tasks identified in Experiment 1 of the GeoPackage in CDB IE.

In the tables below, "Baseline" refers to metrics based on the source ShapeFiles. Option 1A, Option 1C, and Option 1D refer to the sub-options for Experiment 2 (one to one transformation of Shapefiles into GeoPackages).

Table 1. Yemen (4 geocells)

			Baseline	Option 1A	Option 1C	Option 1D	Exper. 3	Exper. 4
Dataset(s)	Feat count	PVF count	time	time	time	time	time	time
tree	64091	440	8	7	7	16	6	2
light	60	13	<1	<1	<1	<1	1	1
cultural	16502	409	12	9	5	7	5	4
powerline	975	20	<1	<1	<1	<1	1	1
railroad	0	0	0	0	0	0	0	0
total time			21	17	13	24	14	8
file count			8224	2056	1023	1023	10	10
size (MB)			34.2	152.5	161.9	165.9	57.6	38.1

Table 2. Downtown Los Angeles (1 geocell)

			Baseline	Option 1A	Option 1C	Option 1D	Exper. 3	Exper. 4
Dataset(s)	Feat count	PVF count	time	time	time	time	time	time
tree	2	1	<1	<1	<1	<1	<1	<1
light	0	0	0	0	0	0	0	0
cultural	1730622	1948	9:01	7:13	3:06	3:34	3:26	3:41
powerline	1208	56	2	1	1	1	<1	1
railroad	1386	4	1	<1	<1	<1	<1	1
total time			9:04	7:15	3:08	3:36	3:27	3:44
file count			12540	4180	2090	2090	4	4
size (MB)			2185.7	2309.2	958.5	1021.5	791.5	798.0

Table 3. Greater Los Angeles (4 geocells)

			Baseline	Option 1A	Option 1C	Option 1D	Exper. 3	Exper. 4
Dataset(s)	Feat count	PVF count	time	time	time	time	time	time
tree	5	2	<1	1	<1	<1	1	<1
light	0	0	0	0	0	0	1	<1
cultural	3138841	6013	15:02	12:02	6:14	7:25	6:57	7:17
powerline	3932	160	1	1	1	1	1	1
railroad	9367	87	1	1	1	1	1	<1
total time			15:04	12:05	6:16	7:27	7:01	7:19
file count			38961	12986	6493	6493	14	14
size (MB)			3738.2	4275.9	1958.6	2067.0	1335.7	1339.3

6.1.4. Legend

- Feat count: feature count of valid features found of the given type
- PVF count: primary vector file count, after validation, for the given type (i.e. only counting .shp files for Experiment 1 or .gpkg files for Experiment 2.)
- Time: in minute:second notation when over 1 minute, else in seconds
- The cultural feature data set is from both 100_GSFeatures (S001_T001 & S002_T001) and 101_GTFeatures (S001_T001)
- File count: total number of files from 100_GSFeatures, 101_GTFeatures, 202_RailroadNetwork & 203_PowerLineNetwork
- Size: storage, in MB, used by all the files from 100_GSFeatures, 101_GTFeatures, 202_RailroadNetwork & 203_PowerLineNetwork

6.1.5. Notes and observations

- All source CDB files were on a local RAID drive so network traffic did not contribute to the timings.
- In the Greater Los Angeles database, there somehow were more features of some types coming from geopackage files compared to shape files (3140180 instead of 3138841 cultural features, and 4012 instead of 3932 powerline features), but there were also over 1000 warnings from OGR during conversion and while reading of the type "Warning 1: Unable to parse srs_id '100000' well-known text "." After the 1000th such warning, also got "More than 1000 errors or warnings have been reported. No more will be reported from now." Perhaps the conversion from .shp to .gpkg with ogr2ogr.exe generated these excess invalid files. These warnings appeared in the Downtown LA database as well, but the feature counts matched after conversion. Checking any further downstream for discrepancies in the processing pipeline was not performed.

- For the powerline network dataset, statistics include both the tower point features and the wire lineal features.
- There's a slight increase in the file size in the Los Angeles databases when comparing the results of Experiment 3 and Experiment 4. However, there is a significant decrease in the size of the Yemen database. From a quick inspection of the data, this seems to correlate with the fact that almost all the cultural features in Los Angeles come from 100_GSFeatures which require unique records per instance, whereas for Yemen the majority of cultural features come from 101_GTFeatures.
- Experiment 3 has slightly better timings for large-count datasets than Experiment 4 in our use case since we scan each LOD in order, so having LODs in separate layers in the option 3 GeoPackage performs better.

6.1.6. Compusult metrics from Experiment 2

Approach: One GeoPackage per LOD per dataset

CDB: CDBYemen_4.0.0

Available Datasets:

- 101_GTFeature
- 100_GSFeature
- 401_Navigation
- 201_RoadNetwork

Number of ShapeFiles processed: 358 Number of GeoPackages created: 18 Total byte size of ShapeFiles (bytes): 3,569,324 Total byte size of GeoPackages (bytes): 41,715,712 Elapsed time (seconds): 173

6.2. FlightSafety Experiment Results

6.2.1. FlightSafety International's Use Case for CDB

FlightSafety has developed both a CDB generation tool and a CDB Publisher client. The performance requirements of the CDB Publisher are much greater than CDB generation, so this report will focus on loading and consuming a CDB dataset. The CDB Publisher uses a CDB data store as the source data for building the synthetic environment for FlightSafety's VITAL 1100 image generator system. These systems are used for pilot training on a variety of flight simulator systems. The Publisher does not do any preprocessing of the CDB dataset; all CDB data that it consumes is discovered and loaded during the publishing. This approach was chosen due to the world-wide scope of CDB and unknown quantity of content. The CDB specification's structure makes it easy to find the file(s) containing the data needed for the synthetic environment creation. Based on the flight training system requirements, an appropriate level of detail of vector and model data is discovered and loaded. The Publisher adapts to the available levels of detail of vector data, and the flight characteristics of the training device. The publishing system is primarily in C, and the testing was all performed with C libraries and code. The Shapefile API that is tested is a custom FlightSafety

library, optimized for faster performance.

6.2.2. FlightSafety Experiment Focus

The experiments were focused on just-in-time uses of CDB, similar to how a FlightSafety visual system would use the data. Statistics were collected on the original CDB dataset, and the converted GeoPackage CDB datasets. These were used to infer the cost of database configuration management and transmission/deployment to a training device. Testing was done on both currently encoded CDB shapefiles, and on converted GeoPackage encoded files (covering options 1, 3, and 4). Tests focused on the latency of loading files, processing data, and closing files. Tests were done on different conversion options and settings to come up with optimal recommendations

6.2.3. FlightSafety Experiment Methodology

This is the methodology used to evaluate, convert and test the CDB datasets using GeoPackage vector encoding.

Data Acquisition:

Three CDB datasets were downloaded (from Cognitics, Presagis, and VATC) and loaded on a system. The datasets were then split into two CDBs, one of which contains all vector data and the other contains everything else. They were linked together using CDB's versioning mechanism, so that the FlightSafety publisher sees the data as a single dataset. Further: - Any official or unofficial extension to the CDB was removed for testing purposes. - Any 0 size vector file was deleted from the CDB with vector data. These were 0 size shp and shx files for datasets that should only be dbf, and cases of 0 size dbt files when they weren't needed alongside their dbf parent file.

Data Evaluation:

All three CDB datasets were flown using FlightSafety's VITAL 1100 image generator and CDB publisher. During the fly-through, any data artifacts were noted and recorded.

Data Conversion:

The python conversion scripts developed by Cognitics, Inc. were downloaded from GitHub. The scripts were modified to properly flatten class-level attributes into the feature table, and to properly handle DBase floating point and logical field types. Index tables were also added to aid SQL queries designed to get back data for a specific CDB vector file. Script changes were published to a public GitHub under a FlightSafety account ([link](#)). When the scripts were run, they created a new output directory for the CDB vector data. The Metadata folder was copied from the original vector CDB version, which then links this GeoPackage version to the rest of the CDB data. The three main conversion scripts used implemented GeoPackage encoding options 1, 3, and 4.

GeoPackage Testing:

The initial data collection centered on the number of vector files and how much disk space was consumed. All full CDB storage devices used a 4kB block size and recorded sizes include the "dead" space due to the minimum block size. The initial tests were testing shapefiles vs. option 1. All vector files were located, and timed on the file open and accessing the data within the file. Total

processing time was recorded and compared between the two encodings. This test accessed the geometry and all the attributes, whether they would have been used by the FlightSafety CDB publisher or not.

The next set of tests involved working with worst case examples and comparing the same file open and access time as before, but for single files. This highlights performance on the largest vector files. The average performance times are reported here.

Further testing was performed to see what the trade offs were between options 1, 3, and 4. These included loading identical vectors (from a single original shapefile) from each of three GeoPackage files converted in different ways:

- GeoPackage Option 1 (Experiment 2 in the IE Activity Plan) was a straight conversion of the shapefile. The GeoPackage contains a single data table with flattened class-level attributes, with the same number of records as the original shapefile
- GeoPackage Option 3 (Experiment 3 in the IE Activity Plan) was a conversion of each CDB dataset's features into a table for each level of detail (LOD) and component selector set, placed into a single GeoPackage (1 per dataset). It also contained the most tables, and typically had more feature records than option 1 but fewer than option 4.
- GeoPackage Option 4 (Experiment 4 in the IE Activity Plan) was a conversion of each CDB dataset's features into a table for each component selector set, placed into a single GeoPackage (1 per dataset). This method placed all levels of detail into the same table, resulting in a handful of tables, but possibly millions of features per table.

Note: *The results for Experiment 1 (Conversion) are provided in the discussions of Option 1, 3, and 4 (IE Experiments 2, 3, and 4).*

6.2.4. FlightSafety Metrics

Original Dataset Statistics

Basic statistics were collected on the original CDB datasets used in the Interoperability Experiment. The CDB storage size and file counts do not include any 0-sized files (they weren't required by the CDB specification) and do not include non-standard extension data. The last two rows represent the proportion of vector data in the CDB, by the percentage of files and storage used. The vector datasets used are: - 100_GSFeature - 101_GTFeature - 102_GeoPolitical - 201_RoadNetwork - 202_RailroadNetwork - 203_PowerlineNetwork - 204_HydrographyNetwork - 401_Navigation

Table 4. Table of Dataset Statistics

	Northwest Pacific	Yemen	Los Angeles
Provider	Cognitics	Presagis	VATC
CDB Geocell Tiles	27	4	1
CDB Storage Size	214 GB	17.4 GB	59.6 GB
CDB File Count	427,536 files	112,837 files	62,895 files
Vector Storage Size	9,152 MB	53.4 MB	2,381 MB
Vector File Count	109,490 files	4714 files	13,075 files

	Northwest Pacific	Yemen	Los Angeles
% of CDB storage as vectors	4.18 %	0.30 %	3.90 %
% of CDB files as vectors	25.6 %	4.18 %	20.8 %

The main takeaway from this table is that vector data does not consume a large amount of storage space, but accounts for a prodigious number of files within a typical CDB. The main driver of file counts are that Shapefiles are a multi-file format, where three (or four with the .prj projection file) files represent a single Shapefile. In addition to the multifile format, CDB uses extra class-level and extended-level attributes encoded as extra DBF files. So anywhere from 3 to 8 files are used to represent a single logical vector file.

Specific Vector File Test Data

Some of the testing below involved loading specific point/linear/areal vectors that represent a single Shapefile. For these tests, examples were found that represent "worst-case" examples of large vector files. These larger files would take more time to load, and most occurred within higher LODs that would lead to larger tables in options 3 and 4. The following table records the specific shapefile data for individual tests.

	Northwest Pacific	Yemen	Los Angeles
Point Vector	N46W124_D101_S002_T001_L04_U15_R12	N12E045_D100_S001_T001_L04_U12_R0	N34W119_D100_S001_T001_L05_U8_R20
Linear Vector	N48W123_D201_S002_T003_L01_U0_R0	N12E045_D201_S002_T003_L00_U0_R0	N34W119_D201_S002_T003_L04_U1_R15
Areal Vector	N47W120_D204_S002_T005_L02_U0_R2	N12E044_D100_S002_T005_L02_U3_R3	N34W119_D204_S002_T005_L03_U4_R7

6.2.5. Shapefile vs. GeoPackage Option 1 (Experiment 2) Testing

Option 1 Conversion Statistics

Before the first set of tests, the CDB datasets were converted one-to-one from shapefiles to GeoPackage, using the option 1 conversion. Dataset statistics were then collected on the new datasets and compared with the original datasets.

	Northwest Pacific	Yemen	Los Angeles
Shapefile Vector Storage Size	9,152 MB	53.4 MB	2,381 MB
Shapefile Vector File Count	109,490 files	4714 files	13,075 files
GeoPackage 1 Storage Size	17,827 MB	157.9 MB	938 MB
GeoPackage 1 File Count	25,083 files	1,146 files	2,615 files

	Northwest Pacific	Yemen	Los Angeles
Relative Size (>1 is larger)	1.95	2.96	0.39
% Fewer Vector Files	77 %	76 %	80 %

File counts for the GeoPackage CDB were between a 4:1 and 5:1 reduction in vector files. The size changes varied dramatically, likely due to how efficient the attributes were packed into the original Shapefile's instance and class-level DBF files. In general, an increase in CDB size is expected using option 1.

Option 1 Testing Focus

The testing focused on the latency of loading and processing the vector data files, and traversing all the geometry features and attributes. This approach was used to simulate a flight simulation client's use of CDB.

Test Procedure 1

The first test was to traverse the entire CDB dataset, find all the vector files and collect the time it took to open, process, and close each vector file. For each dataset, every vector file was located by walking the directory structure, and then the file loading and processing was timed. This test was run 30 times on the smaller CDB datasets (Yemen and Los Angeles) and 10 times on the larger Northwest Pacific dataset. The sum of the file load and process steps are recorded below (while ignoring the file search times).

All Vector Files	Northwest Pacific	Yemen	Los Angeles
Shapefile Timing	835 sec	10.2 sec	27.5 sec
GeoPackage Timing	478 sec	4.2 sec	25.7 sec
GeoPackage Speed Comparison	42% faster	58% faster	6.7% faster
Average Shapefile Storage Size	374 kB	48 kB	923 kB

This table shows, on average, that using GeoPackages are faster than using Shapefiles. These results imply that GeoPackage has a better advantage with smaller files. For example, GeoPackage performed best on Yemen with its relatively small shapefile/vector files. However, there is less of an advantage with larger vector files. Therefore, further testing using larger files is recommended.

Test Procedure 2

The next set of tests focused on some of the largest individual vector files. This test was performed to evaluate some of the worst case examples. The exact file names are mentioned above in the Specific Vector File Test Data section. These test datasets were much larger than the average vector file and cover the three basic geometry types: Points, Line Strings and Polygons. This allowed testing of files that have many attributes compared to coordinates (points), and testing of files with many coordinates compared to the number of attributes (polygons).

- The file size for Shapefiles includes both the instance-level files (.shp, .shx, .dbf) and the class-level attributes (.dbf), but no extended attributes or projection information. The GeoPackage file size was the single .gpkg file.
- The timing numbers include opening the file and traversing the geometry and every attribute in each record, including those that would otherwise not be used by the FlightSafety client. The timing test was performed 100 times alternating between loading from the shapefile CDB dataset, and the equivalent GeoPackage CDB dataset.
- The last row represents the relative performance of GeoPackage as compared to Shapefiles, with a number higher than 1.0 representing increased speed.

Point Vectors	Northwest Pacific	Yemen	Los Angeles
Feature Count	16,384	5,552	4,734
Shapefile Size	1.91 MB	1.40 MB	3.63 MB
GeoPackage Size	3.93 MB	1.46 MB	1.18 MB
Shapefile Read	55.8 ms	64.4 ms	17.4 ms
GeoPackage Read	82.3 ms	36.78 ms	39.9 ms
Relative GeoPackage Performance (>1.0 is faster)	0.678	1.751	0.437

GeoPackage performance numbers were mixed for point data. The GeoPackage performance seems linear with the number of features, but the Shapefile API tested was much faster on one case (Los Angeles) and much slower on another (Yemen).

Note: The Northwest Pacific dataset uses a minimal number of class-level attributes, resulting in a larger flattened GeoPackage size. In contrast, the Los Angeles dataset uses mostly unique class-level attributes, which yields a larger overall Shapefile size, but smaller GeoPackage size because fewer class-level attributes needed to be duplicated.

Line String Vectors	Northwest Pacific	Yemen	Los Angeles
Feature Count	8,183	2,457	3,343
Shapefile Size	1.96 MB	0.71 MB	2.83 MB
GeoPackage Size	2.65 MB	1.08 MB	1.18 MB
Shapefile Read	62.2 ms	26.3 ms	17.4 ms
GeoPackage Read	49.9 ms	19.0 ms	23.1 ms
* Relative GeoPackage Performance (>1.0 is faster)*	1.246	1.383	1.225

The use of GeoPackage increased performance across the board when linear data (22-38%) is processed and used.

PolygonVectors	Northwest Pacific	Yemen	Los Angeles
Feature Count	94	198	127
Shapefile Size	388 kB	387 kB	126 kB
GeoPackage Size	512 kB	556 kB	188 kB
Shapefile Read	9.3 ms	10.0 ms	7.3 ms
GeoPackage Read	6.3 ms	6.4 ms	4.9 ms
Relative GeoPackage Performance (>1.0 is faster)	1.476	1.569	1.502

Larger performance increases for areal data (47% - 56%), at the cost of relatively larger storage size. However, the sample size (number of polygon features) is quite small.

6.2.6. GeoPackage Option 3 and 4 Testing

Please remember that in the FlightSafety presentation of results:

Option 1 = Experiment 2

Option 3 = Experiment 3

Option 4 = Experiment 4

Option 3 & 4 Conversion Statistics

In addition to the one-to-one shapefile to GeoPackage encoding, we wished to also test the other GeoPackage encodings represented by options 3 and 4. Conversions were performed to create these new CDB datasets using the modified python conversion scripts. These were tested against the option 1 CDB datasets used in the previous tests. The dataset statistics (file sizes and counts) are in the table below. Conversion notes include:

- The parts of the file name (dataset code/component selectors/lod/row/column values) were initially stored as strings. Converting these to integers led to about a 10% improvement over the initial string conversion.
- Index tables were created for the parts of the filename that did not comprise the table name. This led to significant improvements that were up to 90% faster than without the index.—Option 3 table names were of the form: "D100_L04_S001_T001". So indexes were created for the row and column values, assuming that a user might want to generate a SQL query on that table's row and column values.—Option 4 table names were of the form: "D100_S001_T001". So indexes were created for the lod, row and column values, assuming that a user might want to generate a SQL query on that table's lod and row and column values.

Conversion Statistics	Northwest Pacific	Yemen	Los Angeles
GeoPackage 1 File Count	25,083 files	1,146 files	2,615 files

Conversion Statistics	Northwest Pacific	Yemen	Los Angeles
GeoPackage 3/4 File Count	161 files	22 files	7 files
% Fewer Vector Files Options 3/4 vs Option 1	99.4 %	98 %	99.7 %

As expected the number of files is much smaller using options 3 or 4.

Conversion Statistics	Northwest Pacific	Yemen	Los Angeles
Shapefile Storage Size	9,152 MB	53.4 MB	2,381 MB
GeoPackage 1 Storage Size	17,827 MB	157.9 MB	938 MB
GeoPackage 3 Storage Size	16,479 MB	59.3 MB	728 MB
GeoPackage 4 Storage Size	16,729 MB	55.5 MB	740 MB
Options 1 & 3 Relative Size (> 1 is larger)	0.92	0.38	0.78
Options 1 & 4 Relative Size (> 1 is larger)	0.94	0.35	0.79
Shapefile vs Option 3 Size (> 1 is larger)	1.80	1.11	0.31
Shapefile vs Option 4 Size (> 1 is larger)	1.83	1.04	0.31

In all cases, the combined GeoPackage datasets required less storage than the option 1 GeoPackage files. This was true even though the combined datasets have index tables that the Option 1 GeoPackages do not have. Note that in all cases even combining the GeoPackage files into a minimal set does not lead to a smaller vector dataset than Shapefiles.

Testing Procedure

The testing focus for comparing the different GeoPackage encoding options was on the latency of loading the GeoPackage file and using SQL queries to return the records converted from a single Shapefile. This approach is similar to the Shapefiles vs GeoPackage testing done above, but the test was constructed slightly differently.

- The GeoPackage was opened and an SQL query was performed to return the data that

represented a single shapefile's vector data. In each query, the number of records in the table varied according to the type of conversion performed, but the amount of data and the values returned by the query were identical.

- The SQL queries used for each GeoPackage option were variations on the following: -- Option 1: `SELECT * FROM 'D100_S001_T001_L04_U12_R0'` -- Option 3: `SELECT * FROM '100_GSFeature_L04_S001_T001' WHERE _UREF='12' AND _RREF='0'` -- Option 4: `SELECT * FROM '100_GSFeature_S001_T001' WHERE _LOD='4' AND _UREF='12' AND _RREF='0'`
- For timing purposes, each GeoPackage's open and query was run 100 times, alternating between each option test in succession.
- The relative speed row shows the performance hit of doing an open on a larger GeoPackage with more tables and more records to search through. For example, a 2.0 represents a test that took twice as long as the option 1 test.

Point Queries	Northwest Pacific	Yemen	Los Angeles
Option 1 Table Size (Count)	16,384	5,552	4,734
Option 1 Read GeoPackage	87.3 ms	39.6 ms	38.6 ms
Option 3 Table Size (Count)	3,375,935	7,766	493,936
Option 3 SQL Query	138.2 ms	59.3 ms	51.5 ms
Speed Relative to Option 1 (>1 is faster)	0.63	0.67	0.75
Option 4 Table Size (Count)	6,865,325	43,122	2,842,150
Option 4 SQL Query	173.9 ms	44.9 ms	79.7 ms
Speed Relative to Option 1 (>1 is faster)	0.50	0.88	0.48

Option 3 GeoPackage file opens are sensitive to the number of tables in the GeoPackage, and tend to dominate the timing of cases with fewer features. Option 4 has fewer tables and faster GeoPackage opens, but is more sensitive to the number of records in the table that need to be searched.

Line String Queries	Northwest Pacific	Yemen	Los Angeles
Option 1 Table Size (Count)	8,183	2,457	3,343
Option 1 Read GeoPackage	53.8 ms	22.3 ms	26.4 ms
Option 3 Table Size (Count)	16,454	2,457	80,697
Option 3 SQL Query	63.6 ms	24.5 ms	28.7 ms

Line String Queries	Northwest Pacific	Yemen	Los Angeles
Speed Relative to Option 1 (>1 is faster)	0.85	0.91	0.92
Option 4 Table Size (Count)	79,512	2,457	149,757
Option 4 SQL Query	52.4 ms	23.0 ms	29.7 ms
Speed Relative to Option 1 (>1 is faster)	1.03	0.97	0.89

In both options 3 and 4, GeoPackage files perform slightly worse, but perform better than the point query because of fewer features returned.

Polygon Queries	Northwest Pacific	Yemen	Los Angeles
Option 1 Table Size (Count)	94	198	127
Option 1 Read GeoPackage	5.4 ms	6.1 ms	4.3 ms
Option 3 Table Size (Count)	207	330	1,238
Option 3 SQL Query	11.3 ms	16.2 ms	6.3 ms
Speed Relative to Option 1 (>1 is faster)	0.48	0.37	0.69
Option 4 Table Size (Count)	2,250	1,531	1,480
Option 4 SQL Query	6.1 ms	7.4 ms	5.0 ms
Speed Relative to Option 1 (>1 is faster)	0.88	0.82	0.86

The overhead of opening GeoPackage files with lots of tables in the option 3 encoding is particularly prominent in the polygon queries. The option 4 encoding is close to the single vector file per GeoPackage timing.

6.2.7. Further GeoPackage Option 3 & 4 Testing

Testing Procedure

The initial SQL query testing only performed a single query per GeoPackage open and close. A more typical use case with options 3 and 4 would be to hold a GeoPackage file open for longer periods of time, and perform more queries per file access. In this test, the same query was performed, but 100 queries were performed while the file was open. There are limitations to the results of this test, as the same query was performed over and over. It was likely that the parts of the file being accessed remained in memory the whole time, and this only measures the time to copy the data out of the

GeoPackage file. But it is a starting point toward understanding the performance of repeated queries in a large file.

The test results show the time per query, plus a 1/100 portion of the GeoPackage open and close time. It also compares this time with Option 1's performance, where there is little gained by keeping the GeoPackage open.

Test Results

Points - 100 Queries	Northwest Pacific	Yemen	Los Angeles
Option 1 - 1 Query	87.3 ms	46.0 ms	38.6 ms
Option 3 - 100 Queries Average	64.6 ms	25.9 ms	26.4 ms
Percent Faster than Option 1	26%	35%	32%
Option 4 - 100 Queries Average	68.2 ms	24.2 ms	23.8 ms
Percent Faster than Option 1	22%	39%	38%

Keeping the GeoPackage open between queries improves performance. But note that not all cases are faster than the original Shapefile performance.

Line Strings - 100 Queries	Northwest Pacific	Yemen	Los Angeles
Option 1 - 1 Query	53.8 ms	22.3 ms	26.4 ms
Option 3 - 100 Queries Average	34.1 ms	11.8 ms	13.3 ms
Percent Faster than Option 1	37%	47%	49%
Option 4 - 100 Queries Average	34.9 ms	11.1 ms	13.3 ms
Percent Faster than Option 1	35%	50%	49%

Polygon - 100 Queries	Northwest Pacific	Yemen	Los Angeles
Option 1 - 1 Query	5.4 ms	6.1 ms	4.3 ms
Option 3 - 100 Queries Average	1.1 ms	1.6 ms	0.78 ms
Percent Faster than Option 1	80%	75%	82%
Option 4 - 100 Queries Average	0.79 ms	1.1 ms	0.85 ms

Polygon - 100 Queries	Northwest Pacific	Yemen	Los Angeles
Percent Faster than Option 1	85%	82%	80%

This approach shows that there is some significant overhead in opening large GeoPackage files. Keeping them open can mitigate some of the overhead. We do not believe that a full client would see this level of performance, but there is a good possibility a client would see improved performance over option 1.

Chapter 7. Guidance

A couple of performance comments (so far):

1. Structure of the data matters. Timing differences in SQL queries on integers rather than strings is enough to matter.
2. As mentioned by others, opening a GeoPackage with lots of tables is slower than having a single table (option 3). .Doing a query to get features out of a very large table is MUCH slower (option 4). I am getting 40x slowdowns for heavily forested areas where I am querying 4700 points out of a table with >2.8M points.
3. The more columns a table has, the larger the slowdown (ie, a query in option 4 vs a query in option3 might take twice as long with 8 columns, but 4 times as long with 30 columns)
 - a. Depending on how much time we have left, testing option 1b might be worthwhile. It should yield faster queries to not flatten class-level attributes into the feature table.

===Holder of info for inclusion

7.1. Who is doing what

7.1.1. Hexagon/Lucid

For our involvement as a participant in the current GeoPackage conversion IE we are at this point planning to complete all listed experiments (1 - 4).

Additionally, we are currently modifying our existing CDB client software to produce the converted vector data in the GeoPackage format and are implementing Option 1d of the conversion strategies.

We will be focused primarily on the client visualization performance for our contribution to the ER. As discussed on last week's call we will also provide some file system metrics after the data conversion. If time permits we will perform the conversions and experiments for all three datasets.

7.2. Metrics captured

7.2.1. FSI

For the client side, FSI documented the following metrics:

- Time to open a GeoPackage (plus SQLite overhead) vs reading/parsing multiple smaller files for shapefile (more I/O operations)
- Time to find/get a layer
- Time to close and dispose of a GeoPackage vs shapefile
- Time to query getting a set of features by SQL query of dataset/lod/row/column vs an rtree SQL search

Chapter 8. Recommendations, Observations, and Conclusions

This section discusses issues, recommendations and change requests.

8.1. Backwards compatibility in the CDB context

8.1.1. FlightSafety observations on backwards compatibility

Our view is that a form of Option 1 is backwards compatible with the existing specification, for the following reasons:

- The CDB concept of a single "logical" file representing a vector file is preserved.
- The CDB conceptual model leans toward mixing the data model and the files on disk. This should be disentangled in a future version of CDB, but it is difficult to do in a minor revision.
- CDB Versioning, where one file from a CDB dataset replaces/hides a file from another CDB dataset, would be difficult to fit together with a GeoPackage that contained "all" the vector files for that dataset. A preferable solution is to evolve the conceptual model to handle this case in CDB 2.0.

8.2. Conclusions

8.2.1. Aechelon

- For the three Experiment 2 options Aechelon tested, the best outcome in both time and file size came from option 1C.
- For Experiments 3 and 4, speed is slightly improved relative to Experiment 3 sub-option 1D but not sub-option 1C. On the other hand, the resulting storage size is markedly improved when compared against all options in Experiment 2, as would be expected. This is because, by design, these Experiments 3 and 4 go against the spirit of CDB data segmentation by file at the LOD level. This makes it more difficult to remove LODs, if so desired, when copying or exporting the CDB vector data. As such, the approaches used in Experiments 3 and 4 may not be as easy to incorporate and adopt as part of the CDB Standard.
- To achieve the improvements in storage while also maintaining the speeds comparable with sub-option 1C and addressing the file-per-LOD issue, Aechelon recommends two additional experiments: (a) where each component selector of each LOD is in its own geopackage file---effectively a variant of sub-option 1C where the U and R references of the same component selectors are combined into one file; and (b) where each dataset's LODs are in a separate geopackage file---effectively a variant of Experiment 3 where instead of storing each LOD in a separate layer in the same geopackage file, each LOD is a separate file.
- If Aechelon were to recommend only one processing alternative, among those in this experiment, for inclusion as an alternate primary vector format in a future OCG CDB revision, it would be option 1C.

8.2.2. FlightSafety

- Shapefiles are a problematic format for several reasons: -- Shapefiles are a multi-file format to represent a single "logical" file. Management of multi-file formats is always a bit harder than handling a single file. -- The CDB specification's use of extra DBF (DBase III) files for class-level and extended-level attribution is not standard for any current GIS tool, and makes the multi-file format problem worse. -- There are other small-level issues with shapefiles, including limited field name length, limited numbers of attributes, limited data types, and a single geometry type per shapefile. -- In all three test CDB datasets, the number of files greatly exceed the proportion of space on disk. The large number of small files causes problems with efficient disk space usage and slows the distribution of a CDB dataset.
- Replacing a single shapefile with a GeoPackage (Option 1) seems to be a better solution for CDB 1.2 or a future CDB 1.x -- Changes to CDB clients involve less work than other options. -- On average, the performance was better with GeoPackage -- On specific worst case tests, Line Strings and Polygons were faster with GeoPackage. Points were mixed, but likely fast enough for this use case. -- It would also fit within CDB's pre-existing versioning mechanism better than the other options, as it can maintain the single "file" paradigm where a vector file replaces/hides another vector file (even if the "file" is really a multi-file format). -- The class-level attributes would be easy to import into a flattened GeoPackage. Importing extended-level attributes as a related table makes sense as well, although redesigning the odd extended attribute encoding would be recommended. -- The SWG is encouraged to implement a single vector encoding per CDB dataset/version, and a way to discover the vector encoding implemented from a file in the Metadata folder.
- GeoPackage options 3 or 4 would be best suited to a future compatibility-breaking version of CDB 2.0. -- CDB client changes would be more extensive, but the performance gain would likely be worth the effort. -- More detail and additional specification requirements would be necessary to describe how to version CDB datasets when that dataset no longer maintains the "traditional" directory and file format that versioning was designed for. -- The SQL query when loading the vectors that would be contained within a single shapefile can take twice as long as when these vectors are in a single GeoPackage, but the open cost can be amortized across many queries if the GeoPackage files remain open for a longer period of time.
- The OGC Interoperability Experiment was a good testing ground to develop familiarity with GeoPackage and to help the SWG come to a better consensus on the use of GeoPackage. We would recommend having more IE's as other parts of the specification are reviewed and improved.

8.3. Recommendations

8.3.1. FlightSafety

GeoPackage Recommendations

- From our experience, GDAL's GeoPackage driver will create non-standard GeoPackage files. The SWG needs to decide if that is allowed, or if a stricter GeoPackage implementation is required. -- GDAL defaults to using "fid" as the primary key in geometry tables. GeoPackage specifies "id" -- As the conversion scripts use GDAL to convert geometry from shapefile to GeoPackage, GDAL

allows the creation of MultiPolygon and MultiLineString features in Polygon and LineString tables (respectively). When this happens, GDAL emits a warning that it is not creating standard GeoPackage files. We would recommend standardizing with GeoPackage, in only allowing a single geometry type within a table. -- The conversion from shapefile logical fields to GeoPackage should be standardized. It would be best to convert any CDB logical field (whether it was logical, string, or integer) into a GeoPackage boolean field. -- The table name should include enough information to be unique, no matter which option is implemented.

- The dataset/component selectors/lod/up/right values must be stored in the table. Integers are recommended for storage for better performance.
- If the SWG decides on using option 3 or 4 for a future version of CDB, then index search tables should be required for better performance when querying data from a specific CDB Tile-LOD.
 1. Forwards compatibility in the CDB context
 2. How to let an application know what vector data encoding is/are available.
 3. Versioning and GPKG and CDB 1.2

Appendix A: Abstract Test Suite

An Abstract Test Suite may be relevant to an Engineering Report.

An Abstract Test Suite is specified in Clause 9 and Annex A of ISO 19105. That Clause and Annex specify the ISO/TC 211 requirements for Abstract Test Suites. Examples of Abstract Test Suites are available in an annex of most ISO 191XX documents, one of the more useful is in ISO 19136. Note that this guidance may be more abstract than needed in an OGC® Implementation Standard.

NOTE

We skip level 2 headers so that asciidoc correctly numbers the subsections in the appendix.

A.1. Test module for conformance level 1

A.1.1. Conformance level 1

Test identifier	/test/case/id
Test purpose:	Confirm that the IUT satisfies all applicable requirements for conformance level 1.
Test method:	Functional testing performed in an automated and/or manual manner. Verify the behaviour of the IUT for the following operations: <ul style="list-style-type: none">• GetCapabilities (mandatory)• DescribeRecord (mandatory)• GetRecords (mandatory)• GetRecordById (mandatory)• GetRepositoryItem (mandatory)• GetDomain (optional)
Requirement:	OGC 07-110: cl. 2.2
Test type:	Capability

A.1.2. Test case for validity of XML response entity

Test identifier	http://www.opengis.net/spec/xxx/conf/WRS.General-ValidResponse
Test purpose:	The XML response entity is valid.
Test method:	Validate content of response entity against corresponding element declaration.
Requirement:	OGC 07-006r1: cl. 10.2.5.1, p. 118
Test type:	Capability

A.2. Test module for conformance level 2

A.2.1. Conformance level 2

Test identifier	/test/case/id
Test purpose:	Confirm that the IUT satisfies all applicable requirements for conformance level 1.
Test method:	Functional testing performed in an automated and/or manual manner. Verify the behaviour of the IUT for the following operations: <ul style="list-style-type: none">• GetCapabilities (mandatory)• DescribeRecord (mandatory)• GetRecords (mandatory)• GetRecordById (mandatory)• GetRepositoryItem (mandatory)• GetDomain (optional)
Requirement:	OGC 07-110: cl. 2.2
Test type:	Capability

A.2.2. Test case for validity of XML response entity

Test identifier	http://www.opengis.net/spec/xxx/conf/WRS.General-ValidResponse
Test purpose:	The XML response entity is valid.
Test method:	Validate content of response entity against corresponding element declaration.
Requirement:	OGC 07-006r1: cl. 10.2.5.1, p. 118
Test type:	Capability

Appendix B: Revision History

Table 5. Revision History

Date	Editor	Release	Primary clauses modified	Descriptions
January 26, 2019	C. Reed	.1	all	Cloned initial version
March 2019	C. Reed	.2	all	Added various overview sections
May 2019	C. Reed	.3	all	Added in experiment results

Appendix C: Bibliography

bibliography::[]